

CS 553 PROGRAMMING ASSIGNMENT-1

EVALUATION DOCUMENT

ROHIT SINGH

A20361198

In this assignment I want to assign a set of jobs (sleep tasks) into a queue which will be done by the workers. The jobs will be varied according to the number of workers and the time for the sleep job. It has to be made sure that the workers do not get duplicate tasks. To do so I do two implementations local and remote. Then I find out the throughput and the efficiency of the implementation.

I have also done the Animoto Clone, which is the remote implementation, with the jobs changes. Instead of sleep jobs the workers need to download HD images and combine them to make a video.

In this assignment I learnt the uses of multithreading through the local implantation and about SQS, DynamoDB and S3 through the remote implementations.

1. AWS Specifications

The specifications used can be summarized as follows:

- Instance : t2.micro
- Region : Oregon(West)
- vCPU : 1
- RAM : 1GB
- OS : Ubuntu Server
- Java : Version 1.7
- SQS (Simple Queuing Service)
- DynamoDB: Read & Write Throughput 160L.
- S3

2. Local

In the local implementation I have implemented multithreading using the **Executor Service** and Queues using the collection class. The experiment reads jobs from the file one by one and pushes it into the queue. Before the tasks can be pushed into the queue, there needs to be some amount of string handling to be done, since, the jobs written in the file are in the following format

```
sleep 0  
sleep 0  
.....  
sleep 0
```

The time to sleep (in the above example is 0ms) is pushed into the queue. A pool of threads are created, and one by one the tasks are popped from the queue and assigned to the threads. The threads keep on sensing the queue until the queue is empty, as shown below:

```
while ( queue.isEmpty() == false )  
{  
    //pop tasks one by one and feed it into the executer service  
    Runnable worker = new WorkerThread(popTask(queue));  
    executor.execute(worker);  
}
```

To execute this experiment the following command can be used:

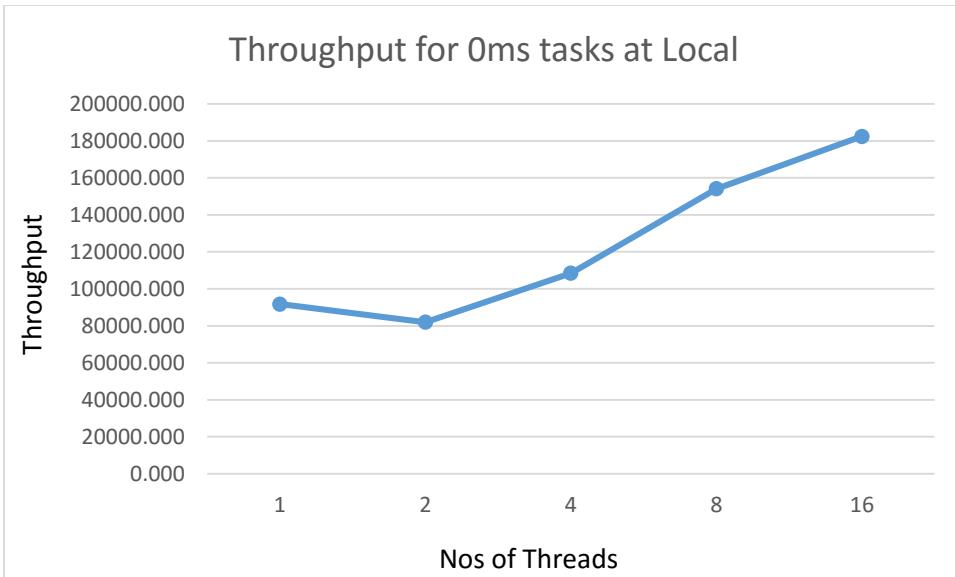
```
java -jar LOCAL.jar <#Threads> <FileName>
```

The local queue guarantees that once an item is popped the item will not get repeated. Hence, the threads will always have unique jobs to perform. Here, the local code emulates the client and the threads as the workers when compared to the remote experiment.

For evaluation we take experiments with 10,000 tasks of 0ms so as to find the throughput of the system.

Table 1 Local experiment of 0ms tasks with 10000 tasks per thread

Nos of Threads	Time taken (in sec)	Throughput (tasks/sec)
1	0.109	91743.119
2	0.122	81967.213
4	0.09225	108401.084
8	0.064875	154142.582
16	0.0548125	182440.137



From the graph above we can observe that the throughput increases with the increase in threads. This is obvious since the number of jobs done parallel will be the proportional to the number of workers. However, it is not always true, as the threads will itself take time to co-ordinat. This becomes an overhead.

To get the efficiency Tasks per worker per instruction we take 3 cases:

- 10ms tasks with 1000 tasks per thread
- 1s tasks with 100 tasks per thread
- 10s tasks with 10 tasks per thread

To calculate the efficiency, I used this formula:

$$\text{Efficiency} = \frac{\text{Time per task} * \text{Tasks per workers}}{\text{Time taken}} * 100 = \frac{\text{Idle Time}}{\text{Time Taken}} * 100$$

Table 2 Local experiment of 10ms tasks with 1000 tasks per thread

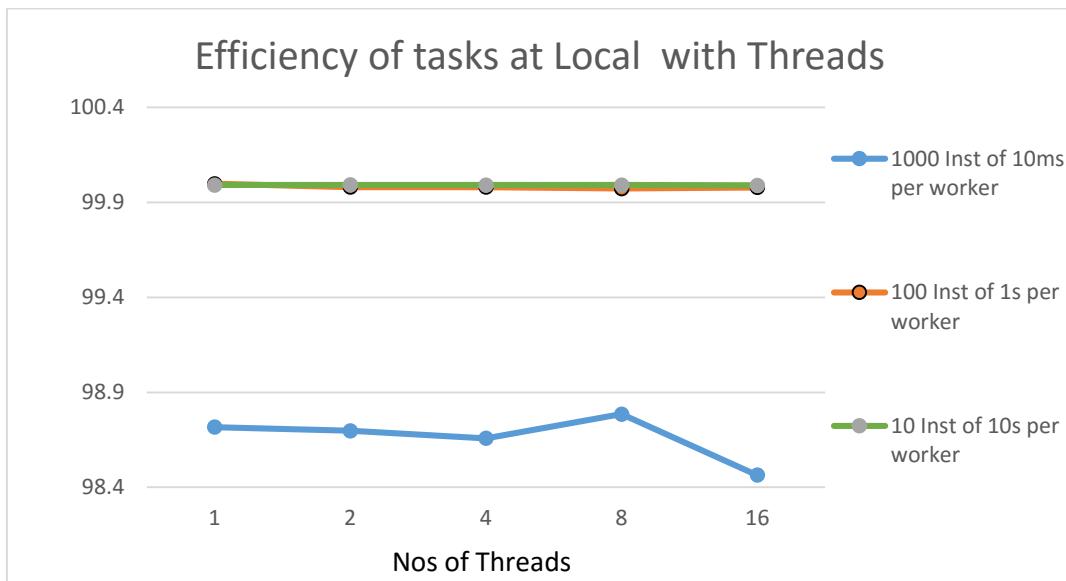
Nos of Threads	Time taken (in sec)	Efficiency(%)
1	10.13	98.7167
2	10.132	98.6972
4	10.136	98.6582
8	10.123	98.7849
16	10.156	98.464

Table 3 Local experiment of 1s tasks with 100 tasks per thread

Nos of Threads	Time taken(in sec)	Efficiency(%)
1	100.002	99.998
2	100.02	99.98
4	100.02	99.98
8	100.027	99.973
16	100.021	99.979

Table 4 Local experiment of 1s tasks with 10 tasks per thread

Nos of Threads	Time taken(in sec)	Efficiency(%)
1	100.009	99.991
2	100.008	99.992
4	100.009	99.991
8	100.009	99.991
16	100.01	99.99



The efficiency for the local is high when the number of instructions are low and high when the number of instructions are high. Since, there is no excess overhead occurring the efficiency is nearly 100%. When the number of instructions are low the probability of the threads not working properly increases.

Screenshots

```
ubuntu@ip-172-31-30-225:~  
wincom@wincom-Studio-XPS-435T:~/Desktop/pa3  
ubuntu@ip-172-31-30-225:~$ ls  
create_sleep_tasks implement jdk-7u1-linux-i586.rpm LOCAL.jar  
----- Tasks of sleep 0 with 10000 instruction (10000 per worker) for 1 workers -----  
Read File  
Queue Created  
Output of queue.isEmpty():false  
Finished all 1 threads in time:109ms  
Time for Execution Full Program for Thread count 1 : .305478186 s  
  
----- Tasks of sleep 0 with 20000 instruction (10000 per worker) for 2 workers -----  
Read File  
Queue Created  
Output of queue.isEmpty():false  
Finished all 2 threads in time:244ms  
Time for Execution Full Program for Thread count 2 : .464798640 s  
  
----- Tasks of sleep 0 with 40000 instruction (10000 per worker) for 4 workers -----  
Read File  
Queue Created  
Output of queue.isEmpty():false  
Finished all 4 threads in time:369ms  
Time for Execution Full Program for Thread count 4 : .579368963 s  
  
----- Tasks of sleep 0 with 80000 instruction (10000 per worker) for 8 workers -----  
Read File  
Queue Created  
Output of queue.isEmpty():false  
Finished all 8 threads in time:519ms  
Time for Execution Full Program for Thread count 8 : .729981611 s  
  
----- Tasks of sleep 0 with 160000 instruction (10000 per worker) for 16 workers -----  
Read File  
Queue Created  
Output of queue.isEmpty():false  
Finished all 16 threads in time:877ms  
Time for Execution Full Program for Thread count 16 : 1.110084864 s  
  
----- Tasks of sleep 10 with 1000 instruction (1000 per worker) for 1 workers -----  
Read File  
Queue Created  
Output of queue.isEmpty():false  
Finished all 1 threads in time:10130ms  
Time for Execution Full Program for Thread count 1 : 10.237437917 s  
  
----- Tasks of sleep 10 with 2000 instruction (1000 per worker) for 2 workers -----  
Read File  
Queue Created  
Output of queue.isEmpty():false
```

```
ubuntu@ip-172-31-30-225:~  
wincom@wincom-Studio-XPS-435T:~/Desktop/pa3  
ubuntu@ip-172-31-30-225:~$  
----- Tasks of sleep 10 with 2000 instruction (1000 per worker) for 2 workers -----  
Read File  
Queue Created  
Output of queue.isEmpty():false  
Finished all 2 threads in time:10132ms  
Time for Execution Full Program for Thread count 2 : 10.252714601 s  
  
----- Tasks of sleep 10 with 4000 instruction (1000 per worker) for 4 workers -----  
Read File  
Queue Created  
Output of queue.isEmpty():false  
Finished all 4 threads in time:10136ms  
Time for Execution Full Program for Thread count 4 : 10.282379199 s  
  
----- Tasks of sleep 10 with 8000 instruction (1000 per worker) for 8 workers -----  
Read File  
Queue Created  
Output of queue.isEmpty():false  
Finished all 8 threads in time:10123ms  
Time for Execution Full Program for Thread count 8 : 10.298448393 s  
  
----- Tasks of sleep 10 with 16000 instruction (1000 per worker) for 16 workers -----  
Read File  
Queue Created  
Output of queue.isEmpty():false  
Finished all 16 threads in time:10156ms  
Time for Execution Full Program for Thread count 16 : 10.365642759 s  
  
----- Tasks of sleep 1000 with 100 instruction (100 per worker) for 1 workers -----  
Read File  
Queue Created  
Output of queue.isEmpty():false  
Finished all 1 threads in time:100022ms  
Time for Execution Full Program for Thread count 1 : 100.117320356 s  
  
----- Tasks of sleep 1000 with 200 instruction (100 per worker) for 2 workers -----  
Read File  
Queue Created  
Output of queue.isEmpty():false  
Finished all 2 threads in time:100020ms  
Time for Execution Full Program for Thread count 2 : 100.114168563 s  
  
----- Tasks of sleep 1000 with 400 instruction (100 per worker) for 4 workers -----  
Read File  
Queue Created  
Output of queue.isEmpty():false  
Finished all 4 threads in time:100020ms  
Time for Execution Full Program for Thread count 4 : 100.118284447 s
```

```

ubuntu@ip-172-31-30-225:~ wincom@wincom-Studio-XPS-435T:~/Desktop/pa3
-----Tasks of sleep 1000 with 800 instruction (100 per worker) for 8 workers -----
Read File
Queue Created
Output of queue.isEmpty():false
Finished all 8 threads in time:100027ms
Time for Execution Full Program for Thread count 8 : 100.138094034 s

-----Tasks of sleep 1000 with 1600 instruction (100 per worker) for 16 workers -----
Read File
Queue Created
Output of queue.isEmpty():false
Finished all 16 threads in time:100021ms
Time for Execution Full Program for Thread count 16 : 100.147456393 s

-----Tasks of sleep 10000 with 10 instruction (10 per worker) for 1 workers -----
Read File
Queue Created
Output of queue.isEmpty():false
Finished all 1 threads in time:100009ms
Time for Execution Full Program for Thread count 1 : 100.102128060 s

-----Tasks of sleep 10000 with 20 instruction (10 per worker) for 2 workers -----
Read File
Queue Created
Output of queue.isEmpty():false
Finished all 2 threads in time:100008ms
Time for Execution Full Program for Thread count 2 : 100.100700593 s

-----Tasks of sleep 10000 with 40 instruction (10 per worker) for 4 workers -----
Read File
Queue Created
Output of queue.isEmpty():false
Finished all 4 threads in time:100009ms
Time for Execution Full Program for Thread count 4 : 100.103966711 s

-----Tasks of sleep 10000 with 80 instruction (10 per worker) for 8 workers -----
Read File
Queue Created
Output of queue.isEmpty():false
Finished all 8 threads in time:100009ms
Time for Execution Full Program for Thread count 8 : 100.104019732 s

-----Tasks of sleep 10000 with 160 instruction (10 per worker) for 16 workers -----
Read File
Queue Created
Output of queue.isEmpty():false
Finished all 16 threads in time:100010ms
Time for Execution Full Program for Thread count 16 : 100.109011735 s

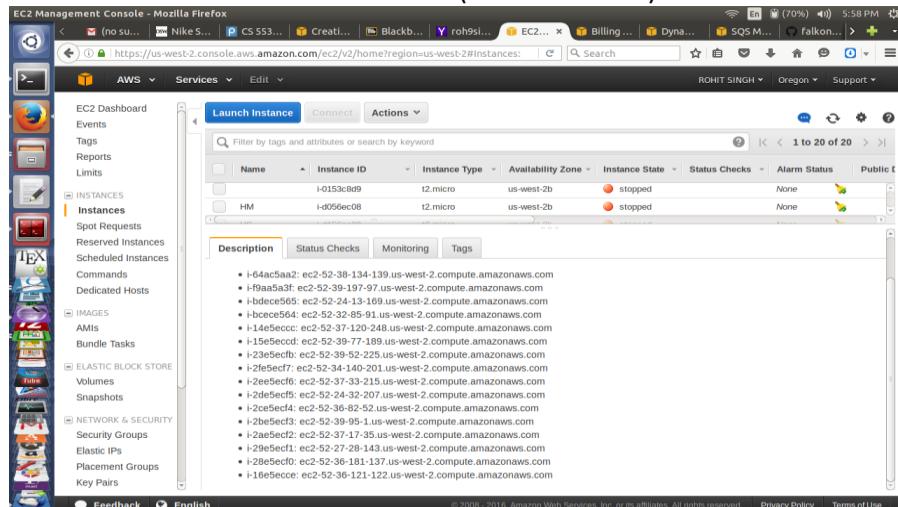
```

3. Remote

In the Remote implementation I have implemented a client and a set of workers which vary from 1, 2, 4, 8 and 16 in number.

The experiment uses SQS and DynamoDB to make a queue and check for duplications respectively.

I have created 17 nodes a client and 16 workers (named 1 to 16) as shown below:



The **client** submits the tasks into a queue here is SQS using a command line interface. After the client submits all the tasks into the queue, it then submits the URL to the queue into DynamoDB into a Table named **QueueURL**, so that the workers have access to the SQS. The client pushes the message as shown below:

```
1 sleep 0
2 sleep 0
.....
10000 sleep 0
```

The first number present in each line provides each task with a unique id, this id can be used later by the workers to keep track of duplicated data. The SQS service is used to handle the queue of requests to load balance across multiple workers. The workers also create a response Queue, so that the client know that the tasks has been done. The following snapshot shows the SQS named **Sleep01** with tasks of 0 ms.

The screenshot shows the AWS SQS Management Console in Mozilla Firefox. The browser address bar shows the URL: <https://console.aws.amazon.com/sqs/home?region=us-west-2#queue-browser:selected>. The main interface displays two SQS queues:

Name	Messages Available	Messages in Flight	Created
ResponseSleep01	1	0	2016-05-03 16:49:28 GMT-05:00
Sleep01	10,000	0	2016-05-03 16:46:25 GMT-05:00

Below the queue list, the details for the **Sleep01** queue are shown:

Details	Permissions	Redrive Policy	Monitoring
Name: Sleep01 URL: https://sqs.us-west-2.amazonaws.com/759689478348/Sleep01 ARN: arn:aws:sqs:us-west-2:759689478348:Sleep01 Created: 2016-05-03 16:46:25 GMT-05:00 Last Updated: 2016-05-03 16:46:25 GMT-05:00 Delivery Delay: 0 seconds	Default Visibility Timeout: 30 seconds Message Retention Period: 4 days Maximum Message Size: 256 KB Receive Message Wait Time: 0 seconds Messages Available (Visible): 10,000 Messages in Flight (Not Visible): 0 Messages Delayed: 0		

The workers retrieve the tasks from the SQS, whose url is present in the DynamoDB with table name **QueueURL**. The workers pull data from the SQS and starts executing them, independently without other workers' intervention. The workers keep on running until the SQS is empty.

The following is the snapshot of the table QueueURL where the Client has submitted the SQS's url for the workers to access the data.

The screenshot shows the AWS DynamoDB console in Mozilla Firefox. The left sidebar lists various AWS services, with 'DynamoDB' selected. Under 'Tables', there is a search bar and a dropdown menu with 'Name' selected. Below it, two items are listed: 'Jobs' and 'QueueURL'. 'QueueURL' is highlighted with a blue selection bar. The main content area is titled 'QueueURL' and shows the 'Items' tab selected. It displays a single item with the key 'name' and the value 'https://sq.s.us-west-2.amazonaws.com/759689478348/Sleep01'. Other tabs like 'Metrics', 'Alarms', 'Capacity', 'Indexes', and 'Triggers' are also visible at the top of the main panel.

Since, SQS cannot guarantee that messages from SQS are delivered exactly once, we use DynamoDB to keep track. I created a new table called **Jobs** which is common to all the workers. The jobs table has 2 attributes: **name** and **task**. The workers before executing the job, checks for the existence of the job in the table. Since, all the tasks are same, the unique id associated with the task can be used to identify the duplication. If a worker finds that the unique id already exists in the table, it will skip the task and pull a new data from the SQS. Else it will insert the task it is going to do and then run the task.

The following is a snapshot of 2 workers working on the same queue, however they do not get duplicate tasks from the queue. Since, the unique id is being monitored by the workers from the jobs table. This also throws light on the fact that all workers do not get equal number of tasks. The following is a simulation of 10s tasks with 10 tasks per worker. So, for 2 workers there are 20 tasks in the queue, out of which the 1st worker has done 11 unique tasks and the 2nd 9 unique tasks. However, it should have been 10 unique tasks each. This gives rise to lesser efficiency, since the 2nd worker will sit idle while the other finishes the extra task.

```

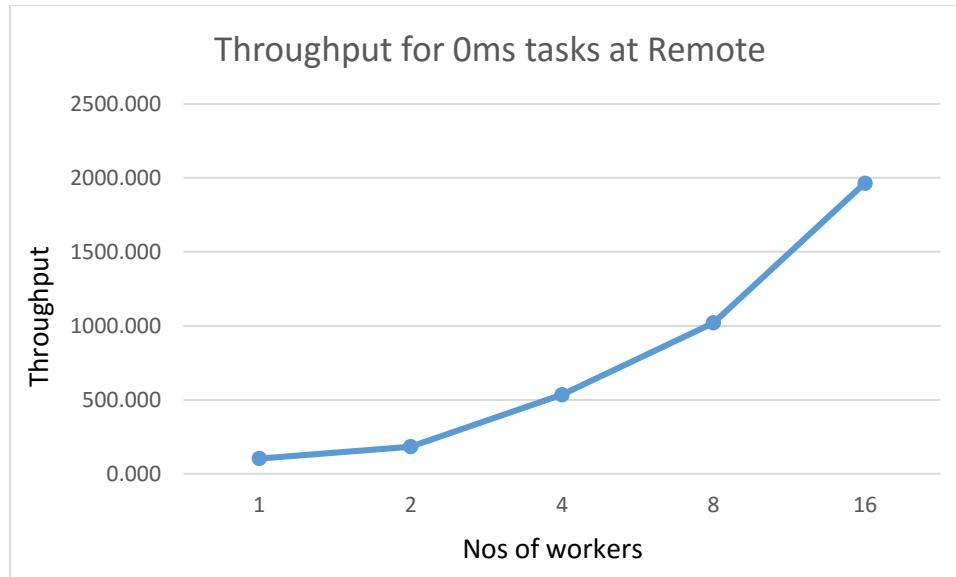
ubuntu@ip-172-31-27-125: ~
At Constructor SimpleQueueServiceSample URL:https://sqs.us-west-2.amazonaws.com/759689478348/Sleep100002
Starting the Jobs for the Worker.....
Reading the Jobs from the Queue and checking for Duplication.....
Unique id:3
Unique id:14
Unique id:1
Unique id:19
Unique id:12
Unique id:7
Unique id:9
Unique id:13
Unique id:20
Unique id:17
Unique id:4
Finished all threads in time:113278
ubuntu@ip-172-31-30-225:~ 151x22
At Constructor SimpleQueueServiceSample URL:https://sqs.us-west-2.amazonaws.com/759689478348/Sleep100002
Starting the Jobs for the Worker.....
Reading the Jobs from the Queue and checking for Duplication.....
Unique id:2
Unique id:10
Unique id:6
Unique id:11
Unique id:15
Unique id:5
Unique id:8
Unique id:16
Unique id:18
Finished all threads in time:114239
ubuntu@ip-172-31-27-125:~ 151x20

```

For evaluation we take experiments with 10,000 tasks of 0ms so as to find the throughput of the system.

Table 5 Remote experiment of 0ms tasks with 10000 tasks per thread

Nos of Threads	Time taken (in sec)	Throughput (tasks/sec)
1	97.02	103.072
2	54.45	183.655
4	18.7	534.759
8	9.8	1020.408
16	5.09	1964.637



From the graph above we can observe that the throughput increases with the increase in workers. This is obvious since the number of jobs done parallel will be the proportional to the number of workers. However, it is not always true, as there might be an uneven distribution of tasks amongst the workers, there might be workers who are sitting idle. Hence, we need to schedule equal or nearly equal amount of jobs to the workers.

To get the efficiency Tasks per worker per instruction we take 3 cases:

- 10ms tasks with 1000 tasks per worker
- 1s tasks with 100 tasks per worker
- 10s tasks with 10 tasks per worker

To calculate the efficiency, I used this formula:

$$\text{Efficiency} = \frac{\text{Time per task} * \text{Tasks per workers}}{\text{Time taken}} * 100 = \frac{\text{Idle Time}}{\text{Time Taken}} * 100$$

Table 6 Remote experiment of 10ms tasks with 1000 tasks per thread

Nos of Workers	Time taken (in sec)	Efficiency(%)
1	32.19	31.0655
2	56.76	35.2361
4	97.89	40.8622
8	181.17	44.1574
16	316.89	50.4907

Table 7 Remote experiment of 1s tasks with 100 tasks per thread

Nos of Workers	Time taken(in sec)	Efficiency(%)
1	155.5	64.3087
2	289.9	68.9893

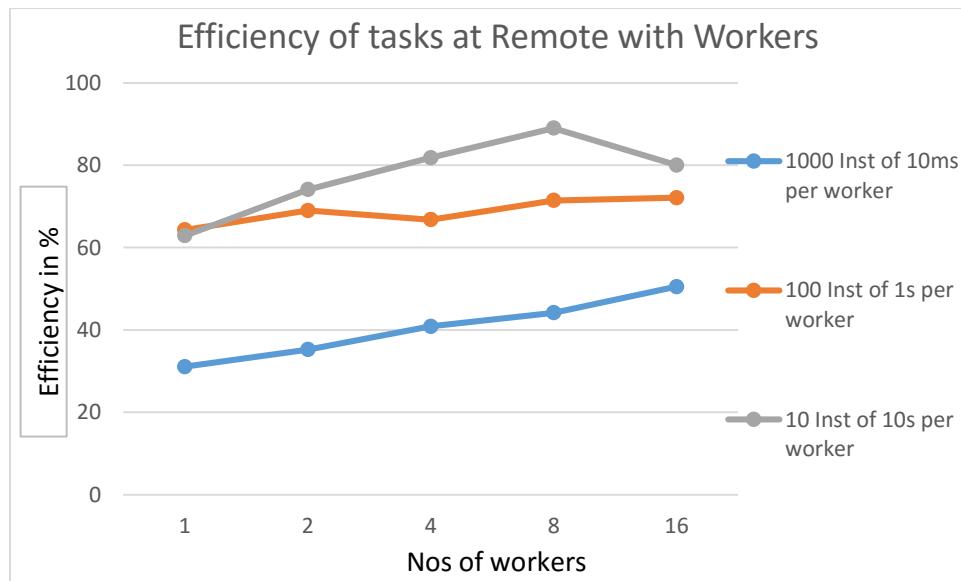
4	598.99	66.7791
8	1119.9	71.4349
16	2219.07	72.1023

Table 8 Remote experiment of 1s tasks with 10 tasks per thread

Nos of Workers	Time taken(in sec)	Efficiency(%)
1	159	62.8931
2	270	74.0741
4	489	81.7996
8	899	88.9878
16	1999	80.04

Since, the workers can be doing random number of tasks at any instance of time, hence I consider the maximum time taken by the workers, as the final time.

It is observed that as we increase the sleep time the time taken by the workers will increase. It is true because the idle time increases.

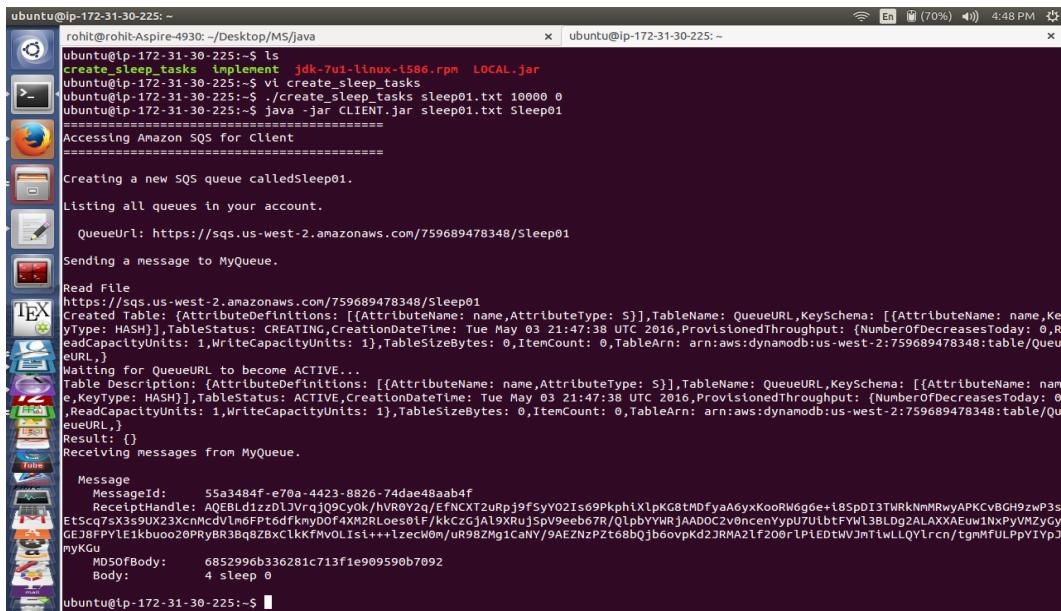


However, if we observe the Efficiency graph in Figure above, we see that the 10ms job with 1000 per worker has very less efficiency compared to the other two jobs. The reason being that there might be few workers who won't be working properly. Some workers might be sitting idle while others doing the maximum tasks. The jobs aren't getting distributed properly amongst the workers. With the increase in workers the number of total tasks also increases (i.e. **Total Task=Worker*Tasks per Worker**). So the time tasks will also increase fairly, as the works aren't distributed properly.

As the sleep task increases the probability that the workers will have tasks evenly distributed will be greater.

The other reason for low efficiency is due to the duplication check. As we keep track of the tasks done by putting it into the DynamoDB. A worker first checks if the task is already present in the table (Jobs). The Jobs is like a Log Book, where each worker reports the task it is going to do. Since, each task, though same, has a unique id associated with it, the id acts as a primary key for the workers to distinguish the unique tasks.

Snapshots



```

ubuntu@lp-172-31-30-225:~ rohit@rohit-Aspire-4930:~/Desktop/M5/java
ubuntu@lp-172-31-30-225:~ ls
create_sleep_tasks implement jdk-7u1-linux-i586.rpm LOCAL.jar
ubuntu@lp-172-31-30-225:~ vi create_sleep_tasks
ubuntu@lp-172-31-30-225:~ ./create_sleep_tasks sleep01.txt 10000 0
ubuntu@lp-172-31-30-225:~ java -jar CLIENT.jar sleep01.txt Sleep01
=====
Accessing Amazon SQS for Client
=====
Creating a new SQS queue calledSleep01.

Listing all queues in your account.

QueueUrl: https://sqs.us-west-2.amazonaws.com/759689478348/Sleep01

Sending a message to MyQueue.

Read File
https://sqs.us-west-2.amazonaws.com/759689478348/Sleep01
Created Table: {AttributeDefinitions: [{AttributeName: name, AttributeType: S}], TableName: QueueURL, KeySchema: [{AttributeName: name, KeyType: HASH}], TableStatus: CREATING, CreationDateTime: Tue May 03 21:47:39 UTC 2016, ProvisionedThroughput: {NumberOfDecreasesToday: 0, ReadCapacityUnits: 1, WriteCapacityUnits: 1}, TableSizeBytes: 0, ItemCount: 0, TableArn: arn:aws:dynamodb:us-west-2:759689478348:table/QueueURL}

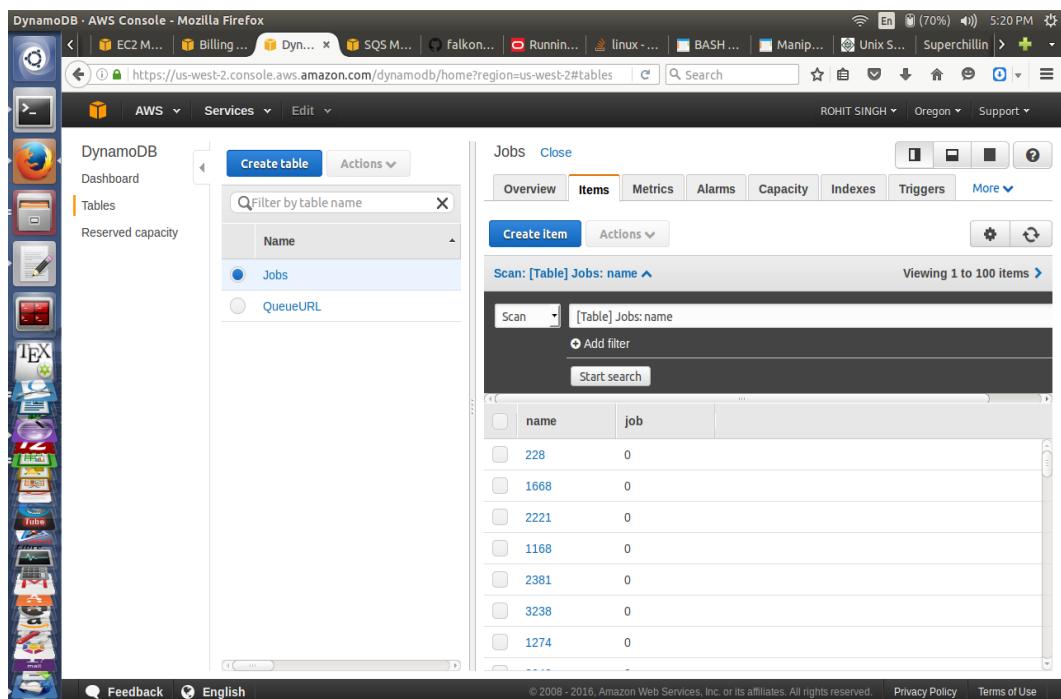
Waiting for QueueURL to become ACTIVE...
Table Description: {AttributeDefinitions: [{AttributeName: name, AttributeType: S}], TableName: QueueURL, KeySchema: [{AttributeName: name, KeyType: HASH}], TableStatus: ACTIVE, CreationDateTime: Tue May 03 21:47:39 UTC 2016, ProvisionedThroughput: {NumberOfDecreasesToday: 0, ReadCapacityUnits: 1, WriteCapacityUnits: 1}, TableSizeBytes: 0, ItemCount: 0, TableArn: arn:aws:dynamodb:us-west-2:759689478348:table/QueueURL}
Result: {}

Receiving messages from MyQueue.

Message
MessageId: 55a3484f-e70a-4423-8826-74dae48aab4f
ReceiptHandle: AQEBLdizzDLVrjqj09cyOk/HvR0Y2g/EfNCXT2uRpj9fSyY02Is69PpkphiXlpKG8tMDfyA6yxKooRW6g6e+i8SpDI3TWRkNmMRwyAPKCVBGH9zwP3sEtScq7sX3s9UX23xcnMcDvIm6fPT6dfkmvd04Xm2RLoest0f/kkCzGjA19XRujSpv9eeb67R/QlpbYYNRJAADOC2v0ncenVyp07ulbtFyMl3BLdg2ALAXXAeuw1NxPvVNzVGyGEJ8FPYLe1kbuo20PRyBR3Bq8ZBxclkkfWolIis1++iZecqB0n/wUR9BZMg1CaNY/9AEZNzPzt68Qjb6ovpKd2JRMa2lf200Rlp1EdtWVmTiwlLQylrcn/tgmMfULPpYIYpJmyGu
MD5OfBody: 6852996b336281c713fe9e09590b7692
Body: 4 sleep 0

ubuntu@lp-172-31-30-225:~ 

```



DynamoDB - AWS Console - Mozilla Firefox

ROHIT SINGH | Oregon | Support

Jobs

name	job
228	0
1668	0
2221	0
1168	0
2381	0
3238	0
1274	0

DynamoDB · AWS Console - Mozilla Firefox

https://us-west-2.console.aws.amazon.com/dynamodb/home?region=us-west-2#tables

ROHIT SINGH · Oregon · Support

Jobs

Items

Metrics

Alarms

Capacity

Indexes

Triggers

More

Create Item

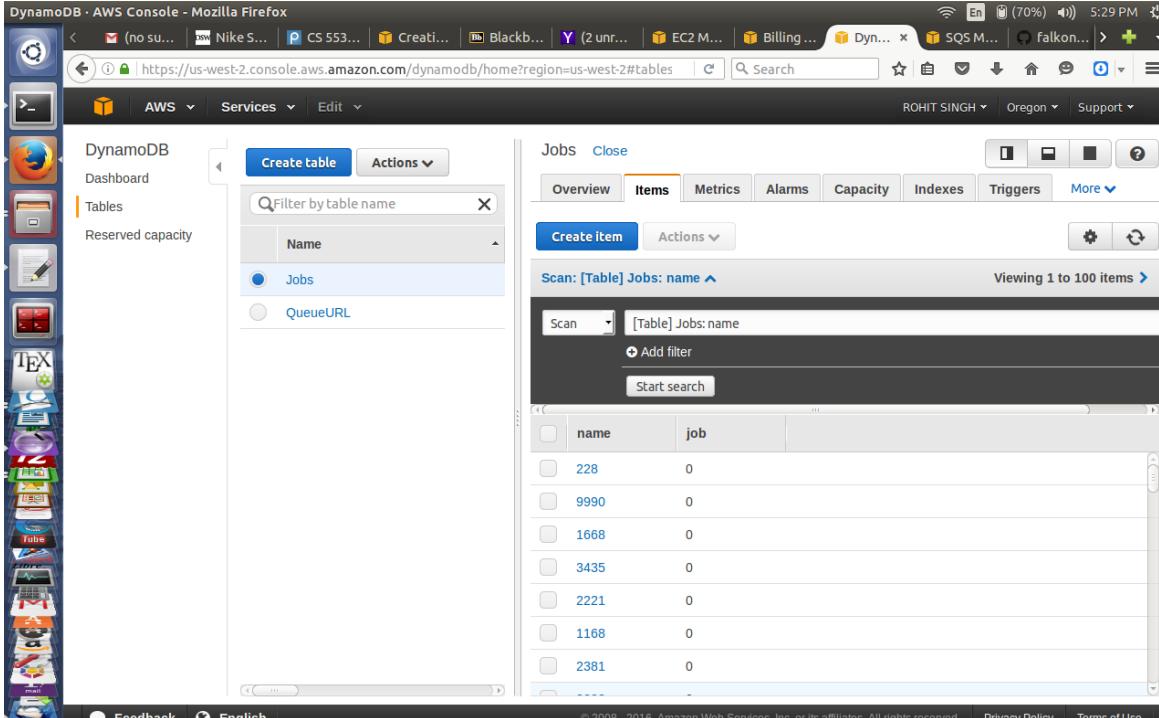
Scan [Table] Jobs: name

Add filter

Start search

	name	job
1	228	0
2	9990	0
3	1668	0
4	3435	0
5	2221	0
6	1168	0
7	2381	0

Feedback English © 2008 - 2016, Amazon Web Services, Inc. or its affiliates. All rights reserved. Privacy Policy Terms of Use



```
ubuntu@ip-172-31-30-225: ~
rohit@rohit-Aspire-4930: ~/Desktop/MS/java      x | ubuntu@ip-172-31-30-225: ~      x | ubuntu@ip-172-31-27-125: ~
ubuntu@ip-172-31-30-225:~$ java -jar WORKER.jar
Table Jobs is already ACTIVE
Access URL:
https://sqS.us-west-2.amazonaws.com/759689478348/Sleep01

Table Description: {AttributeDefinitions: [{AttributeName: name,AttributeType: S}],TableName: Jobs,KeySchema: [{AttributeName: name,KeyType: HASH}],TableStatus: ACTIVE,CreationDateTime: Tue May 03 22:30:45 UTC 2016,ProvisionedThroughput: {NumberOfDecreasesToday: 0,ReadCapacityUnits: 10,WriteCapacityUnits: 10},TableSizeBytes: 0,ItemCount: 0,TableArn: arn:aws:dynamodb:us-west-2:759689478348:table/Jobs,}

At Constructor SimpleQueueServiceSample URL:https://sqS.us-west-2.amazonaws.com/759689478348/Sleep01
Starting the Jobs for the Worker.....

Reading the Jobs from the Queue and checking for Duplication.....

Finished all threads in time:155530
ubuntu@ip-172-31-30-225:~$
```

```

ubuntu@ip-172-31-27-125: ~
[1]:~$ java -jar WORKER.jar
Table Jobs is already ACTIVE
Access URL:
https://sq.s.us-west-2.amazonaws.com/759689478348/Sleep01

Table Description: {AttributeDefinitions: [{AttributeName: name,AttributeType: S}],TableName: Jobs,KeySchema: [{AttributeName: name,KeyType: HASH}],TableStatus: ACTIVE,CreationDateTime: Tue May 03 22:37:21 UTC 2016,ProvisionedThroughput: {NumberOfDecreasesToday: 0,ReadCapacityUnits: 10,WriteCapacityUnits: 10},TableSizeBytes: 0,ItemCount: 0,TableArn: arn:aws:dynamodb:us-west-2:759689478348:table/Jobs,}

At Constructor SimpleQueueServiceSample URL:https://sq.s.us-west-2.amazonaws.com/759689478348/Sleep01
Starting the Jobs for the Worker.....

Reading the Jobs from the Queue and checking for Duplication.....

Finished all threads in time:244312
ubuntu@ip-172-31-30-225:~$ [1]

TEx
[1]:~$ java -jar WORKER.jar 5
Table Jobs is already ACTIVE
Access URL:
https://sq.s.us-west-2.amazonaws.com/759689478348/Sleep01

Table Description: {AttributeDefinitions: [{AttributeName: name,AttributeType: S}],TableName: Jobs,KeySchema: [{AttributeName: name,KeyType: HASH}],TableStatus: ACTIVE,CreationDateTime: Tue May 03 22:37:21 UTC 2016,ProvisionedThroughput: {NumberOfDecreasesToday: 0,ReadCapacityUnits: 10,WriteCapacityUnits: 10},TableSizeBytes: 0,ItemCount: 0,TableArn: arn:aws:dynamodb:us-west-2:759689478348:table/Jobs,}

At Constructor SimpleQueueServiceSample URL:https://sq.s.us-west-2.amazonaws.com/759689478348/Sleep01
Starting the Jobs for the Worker.....

Reading the Jobs from the Queue and checking for Duplication.....

Finished all threads in time:289947
ubuntu@ip-172-31-27-125:~$ [1]

```

```

ubuntu@ip-172-31-33-76: ~
[1]:~$ java -jar WORKER.jar
Table Jobs is already ACTIVE
Access URL:
https://sq.s.us-west-2.amazonaws.com/759689478348/Sleep01

Table Description: {AttributeDefinitions: [{AttributeName: name,AttributeType: S}],TableName: Jobs,KeySchema: [{AttributeName: name,KeyType: HASH}],TableStatus: ACTIVE,CreationDateTime: Tue May 03 22:37:21 UTC 2016,ProvisionedThroughput: {NumberOfDecreasesToday: 0,ReadCapacityUnits: 10,WriteCapacityUnits: 10},TableSizeBytes: 0,ItemCount: 0,TableArn: arn:aws:dynamodb:us-west-2:759689478348:table/Jobs,}

At Constructor SimpleQueueServiceSample URL:https://sq.s.us-west-2.amazonaws.com/759689478348/Sleep01
Starting the Jobs for the Worker.....

Reading the Jobs from the Queue and checking for Duplication.....

Finished all threads in time:544783
ubuntu@ip-172-31-30-225:~$ [1]

TEx
[1]:~$ java -jar WORKER.jar
Table Jobs is already ACTIVE
Access URL:
https://sq.s.us-west-2.amazonaws.com/759689478348/Sleep01

Table Description: {AttributeDefinitions: [{AttributeName: name,AttributeType: S}],TableName: Jobs,KeySchema: [{AttributeName: name,KeyType: HASH}],TableStatus: ACTIVE,CreationDateTime: Tue May 03 22:37:21 UTC 2016,ProvisionedThroughput: {NumberOfDecreasesToday: 0,ReadCapacityUnits: 10,WriteCapacityUnits: 10},TableSizeBytes: 0,ItemCount: 0,TableArn: arn:aws:dynamodb:us-west-2:759689478348:table/Jobs,}

At Constructor SimpleQueueServiceSample URL:https://sq.s.us-west-2.amazonaws.com/759689478348/Sleep01
Starting the Jobs for the Worker.....

Reading the Jobs from the Queue and checking for Duplication.....

Finished all threads in time:598999
ubuntu@ip-172-31-33-77:~$ [1]

[2]:~$ java -jar WORKER.jar
Access URL:
https://sq.s.us-west-2.amazonaws.com/759689478348/Sleep01

Table Description: {AttributeDefinitions: [{AttributeName: name,AttributeType: S}],TableName: Jobs,KeySchema: [{AttributeName: name,KeyType: HASH}],TableStatus: ACTIVE,CreationDateTime: Tue May 03 23:00:37 UTC 2016,ProvisionedThroughput: {NumberOfDecreasesToday: 0,ReadCapacityUnits: 10,WriteCapacityUnits: 10},TableSizeBytes: 0,ItemCount: 0,TableArn: arn:aws:dynamodb:us-west-2:759689478348:table/Jobs,}

At Constructor SimpleQueueServiceSample URL:https://sq.s.us-west-2.amazonaws.com/759689478348/Sleep01
Starting the Jobs for the Worker.....

Reading the Jobs from the Queue and checking for Duplication.....

Finished all threads in time:515257
ubuntu@ip-172-31-27-125:~$ [1]

[2]:~$ java -jar WORKER.jar
Access URL:
https://sq.s.us-west-2.amazonaws.com/759689478348/Sleep01

Table Description: {AttributeDefinitions: [{AttributeName: name,AttributeType: S}],TableName: Jobs,KeySchema: [{AttributeName: name,KeyType: HASH}],TableStatus: ACTIVE,CreationDateTime: Tue May 03 23:00:37 UTC 2016,ProvisionedThroughput: {NumberOfDecreasesToday: 0,ReadCapacityUnits: 10,WriteCapacityUnits: 10},TableSizeBytes: 0,ItemCount: 0,TableArn: arn:aws:dynamodb:us-west-2:759689478348:table/Jobs,}

At Constructor SimpleQueueServiceSample URL:https://sq.s.us-west-2.amazonaws.com/759689478348/Sleep01
Starting the Jobs for the Worker.....

Reading the Jobs from the Queue and checking for Duplication.....

Finished all threads in time:590761
ubuntu@ip-172-31-33-76:~$ [1]

```

```
ubuntu@ip-172-31-30-225: ~
At Constructor SimpleQueueServiceSample URL:https://sns.us-west-2.amazonaws.com/759689478348/Sleep01
Starting the Jobs for the Worker.....
```

Reading the Jobs from the Queue and checking for Duplication.....

```
Finished all threads in time:110135
ubuntu@ip-172-31-30-225: ~
```

Starting the Jobs for the Worker.....

Reading the Jobs from the Queue and checking for Duplication.....

```
Finished all threads in time:110544
ubuntu@ip-172-31-38-144: ~
```

Starting the Jobs for the Worker.....

Reading the Jobs from the Queue and checking for Duplication.....

```
Finished all threads in time:111172
ubuntu@ip-172-31-38-136: ~
```

Starting the Jobs for the Worker.....

Reading the Jobs from the Queue and checking for Duplication.....

```
Finished all threads in time:111990
ubuntu@ip-172-31-27-125: ~
```

Starting the Jobs for the Worker.....

Reading the Jobs from the Queue and checking for Duplication.....

```
Finished all threads in time:111413
ubuntu@ip-172-31-33-76: ~
```

Starting the Jobs for the Worker.....

Reading the Jobs from the Queue and checking for Duplication.....

```
ubuntu@ip-172-31-38-135: ~
Starting the Jobs for the Worker.....
```

Reading the Jobs from the Queue and checking for Duplication.....

```
Finished all threads in time:111404
ubuntu@ip-172-31-38-135: ~
```

Starting the Jobs for the Worker.....

Reading the Jobs from the Queue and checking for Duplication.....

```
Finished all threads in time:110351
ubuntu@ip-172-31-38-146: ~
```

DynamoDB · AWS Console - Mozilla Firefox

https://us-west-2.console.aws.amazon.com/dynamodb/home?region=us-west-2#tables

ROHIT SINGH · Oregon · Support

DynamoDB · AWS · Services · Edit

Dashboard · Tables · Reserved capacity

Create table · Actions

Filter by table name · Name

Jobs · QueueURL

Jobs · Close

Overview · Items · Metrics · Alarms · Capacity · Indexes · Triggers · More

Create item · Actions

Scan: [Table] Jobs: name · Viewing 1 to 100 items

Scan · [Table] Jobs: name · Add filter · Start search

	name	job
228	0	
9990	0	
1668	0	
3435	0	
2221	0	
1168	0	
2381	0	

Feedback · English · 2008 - 2016, Amazon Web Services, Inc. or its affiliates. All rights reserved. · Privacy Policy · Terms of Use

	name	job
228	0	
9990	0	
1668	0	
3435	0	
2221	0	
1168	0	
2381	0	
3238	0	
1274	0	

DynamoDB · AWS Console - Mozilla Firefox

https://us-west-2.console.aws.amazon.com/dynamodb/home?region=us-west-2#tables

ROHIT SINGH · Oregon · Support

DynamoDB · AWS · Services · Edit

Dashboard · Tables · Reserved capacity

Create table · Actions

Filter by table name · Name

Jobs · QueueURL

Jobs · Close

Overview · Items · Metrics · Alarms · Capacity · Indexes · Triggers · More

Create item · Actions

Scan: [Table] Jobs: name · Viewing 1 to 100 items

Scan · [Table] Jobs: name · Add filter · Start search

	name	job
228	0	
1668	0	
2221	0	
1168	0	
2381	0	
3238	0	
1274	0	

Feedback · English · 2008 - 2016, Amazon Web Services, Inc. or its affiliates. All rights reserved. · Privacy Policy · Terms of Use

	name	job
228	0	
9990	0	
1668	0	
3435	0	
2221	0	
1168	0	
2381	0	
3238	0	
1274	0	

DynamoDB - AWS Console - Mozilla Firefox

https://us-west-2.console.aws.amazon.com/dynamodb/home?region=us-west-2#tables

ROHIT SINGH | ROHIT SINGH | Support

DynamoDB

Dashboard

Tables

Reserved capacity

Create table Actions

Filter by table name

Name

Jobs QueueURL

Jobs Close

Overview Items Metrics Alarms Capacity Indexes Triggers More

create item Actions

Scan [Table] Jobs: name Viewing 1 to 100 items

Scan Add filter Start search

	name	job
99	1000	
63	1000	
92	1000	
64	1000	
36	1000	
49	1000	
33	1000	
...	...	

Feedback English

© 2008 - 2016, Amazon Web Services, Inc. or its affiliates. All rights reserved. Privacy Policy Terms of Use

DynamoDB - AWS Console - Mozilla Firefox

https://us-west-2.console.aws.amazon.com/dynamodb/home?region=us-west-2#tables

ROHIT SINGH | ROHIT SINGH | Support

DynamoDB

Dashboard

Tables

Reserved capacity

Create table Actions

Filter by table name

Name

Jobs QueueURL

Jobs Close

Overview Items Metrics Alarms Capacity Indexes Triggers More

create item Actions

Scan [Table] Jobs: name Viewing 1 to 20 items

Scan Add filter Start search

	name	job
18	10000	
16	10000	
2	10000	
13	10000	
8	10000	
9	10000	
1	10000	
...	...	

Feedback English

© 2008 - 2016, Amazon Web Services, Inc. or its affiliates. All rights reserved. Privacy Policy Terms of Use

DynamoDB - AWS Console - Mozilla Firefox

https://us-west-2.console.aws.amazon.com/dynamodb/home?region=us-west-2#tables

ROHIT SINGH | ROHIT SINGH | Support

DynamoDB

Dashboard

Tables

Reserved capacity

Create table Actions

Filter by table name

Name

Jobs QueueURL

Jobs Close

Overview Items Metrics Alarms Capacity Indexes Triggers More

create item Actions

Scan [Table] Jobs: name Viewing 1 to 100 items

Scan Add filter Start search

	name	job
99	1000	
63	1000	
92	1000	
64	1000	
36	1000	
49	1000	
33	1000	
...	...	

Feedback English

© 2008 - 2016, Amazon Web Services, Inc. or its affiliates. All rights reserved. Privacy Policy Terms of Use

4. Animoto Clone [Extra Credit]

Animoto is a cloud-based video creation service that produces video from photos, video clips, and music into video slideshows [7]. In this assignment we need to make a clone of the same using the cloud services provided by AWS namely SQS and S3. For this experiment I have used the **ffmpeg** to create a video out of the photos which have been downloaded from the urls provided by the file as an input, these set of urls will be sent by the Client for the S3 to download using **wget**. This work is to be done by the workers and after each work is done by a worker it submit the video into S3, and returns the link to the S3 to a response queue (SQS) for the client to view. There are about 160 such tasks which is to be done by a set of workers. The worker set varies as 1,2,4,8 and 16.

FFmpeg is a free software project that produces libraries and programs for handling multimedia data. FFmpeg includes libavcodec, an audio/video codec library used by several other projects, libavformat, an audio/video container mux and de-mux library, and the ffmpeg command line program for transcoding multimedia files [8-12]. The following code combines all the images named as image%d.jpg, where %d is a number present in the source path folder.

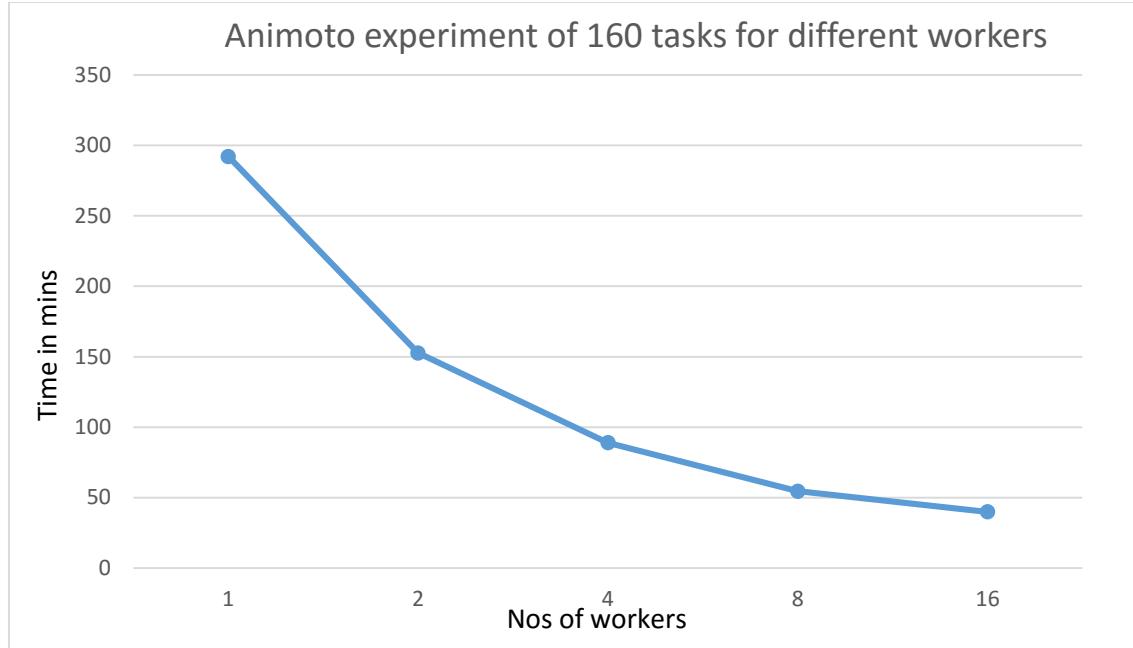
```
ffmpeg -f image2 -i <source path>/image%d.jpg <destination path>/output.mpg
```

All the images are of 1920x1080 resolution. The urls provided by the Client are all 1920x1080 resolution, since we want to make a 60 seconds long video clip with HD 1080P resolution. To make the video 60 sec long I have slowed the fps (frames per second), by using the following code:

```
ffmpeg -i output.mpg -filter:v "setpts=10*PTS" output.mpg
```

Table 9 Animoto experiment of 160 tasks with 60 images

Nos of Threads	Time taken(in min)
1	292.13
2	152.69
4	88.98
8	54.51
16	39.91



The experiment was performed with 160 such tasks with 1 to 16 workers. The time decreases with increasing number of workers as shown in Figure , since the tasks get distributed. The major time required to do the task is to download the HD quality pictures. So, the timing is largely governed by the network bandwidth at that instance.

Snapshots

```
ubuntu@ip-172-31-27-125:~$ java -jar Animoto.jar
ubuntu@ip-172-31-30-225:~ 151x21
Table Jobs is already ACTIVE
Access URL:
https://sqs.us-west-2.amazonaws.com/759689478348/Sleep01
Table Description: {AttributeDefinitions: [{AttributeName: name,AttributeType: S}],TableName: Jobs,KeySchema: [{AttributeName: name,KeyType: HASH}],TableStatus: ACTIVE,CreatedAt:Tue May 03 22:37:21 UTC 2016,ProvisionedThroughput: {NumberOfDecreasesToday: 0,ReadCapacityUnits: 10,WriteCapacityUnits: 10},TableSizeBytes: 0,ItemCount: 0,TableArn: arn:aws:dynamodb:us-west-2:759689478348:table/Jobs,}
At Constructor SimpleQueueServiceSample URL:https://sqs.us-west-2.amazonaws.com/759689478348/Sleep01
Starting the Jobs for the Worker.....
Reading the Jobs from the Queue and checking for Duplication.....
Finished all threads in time:150.22
ubuntu@ip-172-31-30-225:~ [tex]

ubuntu@ip-172-31-27-125:~$ java -jar Animoto.jar
ubuntu@ip-172-31-27-125:~ 151x21
Table Jobs is already ACTIVE
Access URL:
https://sqs.us-west-2.amazonaws.com/759689478348/Sleep01
Table Description: {AttributeDefinitions: [{AttributeName: name,AttributeType: S}],TableName: Jobs,KeySchema: [{AttributeName: name,KeyType: HASH}],TableStatus: ACTIVE,CreatedAt:Tue May 03 22:37:21 UTC 2016,ProvisionedThroughput: {NumberOfDecreasesToday: 0,ReadCapacityUnits: 10,WriteCapacityUnits: 10},TableSizeBytes: 0,ItemCount: 0,TableArn: arn:aws:dynamodb:us-west-2:759689478348:table/Jobs,}
At Constructor SimpleQueueServiceSample URL:https://sqs.us-west-2.amazonaws.com/759689478348/Sleep01
Starting the Jobs for the Worker.....
Reading the Jobs from the Queue and checking for Duplication.....
Finished all threads in time:152.69
ubuntu@ip-172-31-27-125:~ [tex]
```

Object: W1output.mp4

Bucket: animotourl
Name: W1output.mp4
Link: https://s3-us-west-2.amazonaws.com/animotourl/W1output.mp4
Size: 1531904
Last Modified: Wed May 04 18:55:21 GMT-05:00 2016
Owner: rsingh45
ETag: 19735b0ea093a9952ee97151b068c53
Expiry Date: None
Expiration Rule: N/A

Details

Permissions

Metadata

Delete	Body	More Details	Size	Sent	Receive Count
<input type="checkbox"/>	https://s3-us-west-2.amazonaws.com/animotourl/W8output2.mp4	More Details	59 bytes	2016-05-04 19:08:43 GMT-05:00	1
<input type="checkbox"/>	https://s3-us-west-2.amazonaws.com/animotourl/W8output3.mp4	More Details	59 bytes	2016-05-04 19:08:55 GMT-05:00	1
<input type="checkbox"/>	https://s3-us-west-2.amazonaws.com/animotourl/W1output.mp4	More Details	58 bytes	2016-05-04 18:57:55 GMT-05:00	5
<input type="checkbox"/>	https://s3-us-west-2.amazonaws.com/animotourl/W2output.mp4	More Details	58 bytes	2016-05-04 18:57:59 GMT-05:00	5
<input type="checkbox"/>	https://s3-us-west-2.amazonaws.com/animotourl/W10output3.mp4	More Details	60 bytes	2016-05-04 19:09:06 GMT-05:00	1
<input type="checkbox"/>	https://s3-us-west-2.amazonaws.com/animotourl/W2output2.mp4	More Details	59 bytes	2016-05-04 19:07:27 GMT-05:00	2
<input type="checkbox"/>	https://s3-us-west-2.amazonaws.com/animotourl/W3output.mp4	More Details	58 bytes	2016-05-04 19:07:48 GMT-05:00	2
<input type="checkbox"/>	https://s3-us-west-2.amazonaws.com/animotourl/W1output.mp4	More Details	58 bytes	2016-05-04 18:56:14 GMT-05:00	4

100%

Stopped after polling the queue at 0.5 receives/second for 20.2 seconds. Messages shown above are now available to other consumers.

Start Polling for Messages **Stop Now**

Close **Delete Messages**

5. Reference

- [1] http://docs.aws.amazon.com/AWSToolkitEclipse/latest/ug/tke_setup_install.html
- [2] <http://docs.aws.amazon.com/amazondynamodb/latest/developerguide/AppendixSampleDataCodeJava.html>
- [3] <http://docs.aws.amazon.com/amazondynamodb/latest/developerguide/ScanJavaDocumentAPI.html>
- [4] <http://docs.aws.amazon.com/amazondynamodb/latest/developerguide/Expressions.SpecifyingConditions.html>
<http://docs.aws.amazon.com/amazondynamodb/latest/developerguide/Expressions.SpecifyingConditions.html>
- [5] <http://www.avajava.com/tutorials/lessons/how-do-i-save-an-image-from-a-url-to-a-file.html>
- [6] <http://askubuntu.com/questions/699502/ffmpeg-command-not-found>
- [7] <https://en.wikipedia.org/wiki/Animoto>
- [8] <https://en.wikipedia.org/wiki/FFmpeg>
- [9] http://www.ffmpeg.org/faq.html#How-do-I-encode-single-pictures-into-movies_003f
- [10] <https://trac.ffmpeg.org/wiki/How%20to%20speed%20up%20/%20slow%20down%20a%20video>
- [11] <http://stackoverflow.com/questions/24961127/ffmpeg-create-video-from-images>
- [12] <http://www.mkyong.com/java/how-to-execute-shell-command-from-java/>
- [13] http://docs.aws.amazon.com/AWSSimpleQueueService/latest/APIReference/API_CreateQueue.html
- [14] http://docs.aws.amazon.com/AWSSimpleQueueService/latest/APIReference/API_GetQueueUrl.html
- [15] http://docs.aws.amazon.com/AWSSimpleQueueService/latest/APIReference/API_ReceiveMessage.html
- [16] http://docs.aws.amazon.com/amazondynamodb/latest/APIReference/API_GetItem.html
- [17] http://docs.aws.amazon.com/amazondynamodb/latest/APIReference/API_Scan.html
- [18] <http://docs.aws.amazon.com/AmazonS3/latest/API>Welcome.html>